

[udaylunawat@gmail.com \(mailto:udaylunawat@gmail.com\)](mailto:udaylunawat@gmail.com)

```
In [1]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
import matplotlib.pyplot as plt
```

```
In [2]: X, y = make_classification(n_samples=50000, n_features=15, n_informative=10, n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7, random_state=15)
```

```
In [3]: X.shape, y.shape
```

```
Out[3]: ((50000, 15), (50000,))
```

```
In [4]: from sklearn.model_selection import train_test_split
```

```
In [5]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=15)
```

```
In [6]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[6]: ((37500, 15), (37500,)), (12500, 15), (12500,))
```

```
In [7]: from sklearn import linear_model
```

```
In [8]: # alpha : float
        # Constant that multiplies the regularization term.

        # eta0 : double
        # The initial learning rate for the 'constant', 'invscaling' or 'adaptive' schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log', random_state=1, penalty='l2', tol=1e-3, verbose=
clf
```

```
Out[8]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                    early_stopping=False, epsilon=0.1, eta0=0.0001,
                    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                    penalty='l2', power_t=0.5, random_state=1, shuffle=True,
                    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [9]: clf.fit(X=X_train, y=y_train)
```

```
-- Epoch 1
Norm: 0.76, NNZs: 15, Bias: -0.314219, T: 37500, Avg. loss: 0.456332
Total training time: 0.01 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.471117, T: 75000, Avg. loss: 0.394828
Total training time: 0.02 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.581147, T: 112500, Avg. loss: 0.385666
Total training time: 0.03 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.660030, T: 150000, Avg. loss: 0.382073
Total training time: 0.04 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719106, T: 187500, Avg. loss: 0.380409
Total training time: 0.04 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763106, T: 225000, Avg. loss: 0.379544
Total training time: 0.05 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.794459, T: 262500, Avg. loss: 0.379092
Total training time: 0.06 seconds.
-- Epoch 8
Norm: 1.07, NNZs: 15, Bias: -0.819748, T: 300000, Avg. loss: 0.378926
Total training time: 0.07 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837385, T: 337500, Avg. loss: 0.378727
Total training time: 0.08 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.851748, T: 375000, Avg. loss: 0.378574
Total training time: 0.08 seconds.
Convergence after 10 epochs took 0.08 seconds
```

```
Out[9]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0001,
    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
    penalty='l2', power_t=0.5, random_state=1, shuffle=True,
    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [10]: clf.coef_, clf.coef_.shape, clf.intercept_
```

```
Out[10]: (array([[ -0.4167019 ,  0.18364588, -0.14293025,  0.34812889, -0.21566194,
                  0.56418028, -0.44306199, -0.09788853,  0.21219686,  0.17592993,
                  0.1925475 ,  0.0017679 , -0.08466201,  0.34080004,  0.03251288]]),
          (1, 15),
          array([ -0.85174791]))
```

```
In [11]: clf.intercept_[0]
```

```
Out[11]: -0.8517479126246573
```

Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

Instructions

- Load the datasets(train and test) into the respective arrays
- Initialize the weight_vector and intercept term randomly
- Calculate the initial log loss for the train and test data with the current weight and intercept and store it in a list
- for each epoch:
 - for each batch of data points in train: (keep batch size=1)
 - calculate the gradient of loss function w.r.t each weight in weight vector
 - Calculate the gradient of the intercept [check this \(https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing\)](https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing).
 - Update weights and intercept (check the equation number 32 in the above mentioned [pdf \(https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing\)](https://drive.google.com/file/d/1nQ08-XY4zvOLzRX-IGf8EYB5arb7-m1H/view?usp=sharing)):

$$w^{(t+1)} \leftarrow (1 - \frac{\alpha\lambda}{N})w^{(t)} + \alpha x_n (y_n - \sigma((w^{(t)})^T x_n + b^t))$$

$$b^{(t+1)} \leftarrow (1 - \frac{\alpha\lambda}{N})b^{(t)} + \alpha (y_n - \sigma((w^{(t)})^T x_n + b^t))$$

- calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
 - And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
 - append this loss in the list (this will be used to see how loss is changing for each epoch after the training is over)
- Plot the train and test loss i.e on x-axis the epoch number, and on y-axis the loss
- **GOAL:** compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

```
In [236]: w = np.zeros_like(X_train[0])
b = 0
eta0 = 0.0001
alpha = 0.0001
N = len(X_train)
```

```
In [163]: def sigmoid(w,x,b):
return 1/(1+np.exp(-(np.dot(x,w)+b)))
```

```
In [164]: def predict(w,b,x):
pred = []
for i in range(len(x)):
pred.append(sig(w, X[i], b))
return np.array(pred)
```

```
In [165]: import math
def compute_log_loss(true,pred):
    loss = 0
    for t,p in zip(true,pred):
        y1 = t
        y2 = p
        loss += (y1* np.log(y2)) + ((1-y1)* np.log(1-y2))
        # loss = float("{0:.4f}".format(loss))
    return (-1*(loss)/len(true))
```

$$w^{(t+1)} \leftarrow (1 - \frac{\alpha\lambda}{N})w^{(t)} + \alpha x_n (y_n - \sigma((w^{(t)})^T x_n + b^t))$$

$$b^{(t+1)} \leftarrow (1 - \frac{\alpha\lambda}{N})b^{(t)} + \alpha (y_n - \sigma((w^{(t)})^T x_n + b^t)) "$$

```

In [237]: from tqdm import tqdm
a = alpha
l = eta0
train_loss = []
test_loss = []
print("epoch\t log loss")

for epoch in range(30):
    for j in range(N):
        r = np.random.randint(N)
        Xn = X_train[r]
        yn = y_train[r]

        weight_update = ( 1- (a*l)/N) *w + (a*(Xn*(yn-sigmoid(w,Xn,b))))
        intercept_update = (1 - (a*l)/N) *b + (a*(yn-sigmoid(w,Xn,b)))

        w = weight_update
        b = intercept_update

    y_train_pred = map(lambda i: sigmoid(w,i,b), X_train)
    y_test_pred = map(lambda i: sigmoid(w,i,b), X_test)

    loss = compute_log_loss(y_train,y_train_pred)

    train_loss.append(loss)
    test_loss.append(compute_log_loss(y_test,y_test_pred))

    print(epoch, '\t', loss)


    if abs(b-clf.intercept_)[0] < 0.01 and np.allclose(w,clf.coef_,rtol = 1e-01, atol = 1e-02):
        print("converged")
        break

```

epoch	log loss
0	0.4043300564834753
1	0.3884187924554347
2	0.38344655505551983
3	0.3807626284133473
4	0.3795749054664712
5	0.37897757910248303

```
6      0.3787680772563091
7      0.378490169081829
8      0.3783933590871815
9      0.379236863992006
converged
```

```
In [213]: plt.plot(test_loss,label='test_log_loss')
plt.plot(train_loss,label='train_log_loss')
plt.legend()
plt.xlabel('epochs')
plt.ylabel('logloss')
plt.show();
```



```
In [238]: print(w,b)

[-0.42675708  0.18378038 -0.13647368  0.33474876 -0.21133386  0.56084653
 -0.43580897 -0.09492469  0.23051905  0.17307646  0.18122995  0.00582301
 -0.07340738  0.33828899  0.02947331] -0.8485635519484128
```



```
In [215]: def pred(w,b,X):
            N = len(X)
            predict = []
            for i in range(N):
                if sigmoid(w, X[i], b) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
                    predict.append(1)
                else:
                    predict.append(0)
            return np.array(predict)
print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))
```

```
0.9552266666666667
0.95336
```

```
In [216]: print(1-np.sum(y_train - pred(clf.coef_[0], clf.intercept_,X_train))/len(X_train))
print(1-np.sum(y_test - pred(clf.coef_[0], clf.intercept_,X_test))/len(X_test))
```

```
0.95288
0.95256
```

```
In [217]: print(w-clf.coef_)
print(b-clf.intercept_)
np.allclose(w,clf.coef_,rtol = 1e-02, atol = 1e-02)
```

```
[[-0.00353859  0.00196883 -0.00089057 -0.00210887  0.00644648 -0.00087634
 -0.00680417  0.00294099  0.00227197  0.00071759 -0.00043365  0.00125569
  0.00174226  0.0014      -0.00426724]]
[0.00861239]
```

Out[217]: True

In []:

